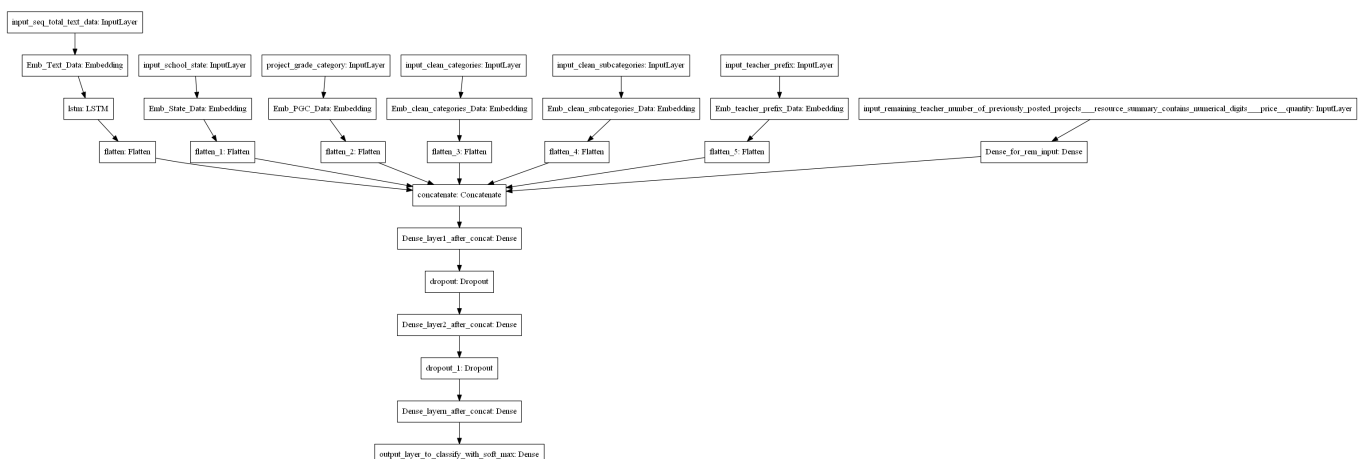


Assignment : 14

1. Preprocess all the Data we have in DonorsChoose [Dataset \(https://drive.google.com/drive/folders/1MIwK7BQMeV8f5CbDDVNLPaFGB32pFN60\)](https://drive.google.com/drive/folders/1MIwK7BQMeV8f5CbDDVNLPaFGB32pFN60) use train.csv
2. Combine 4 essay's into one column named - 'preprocessed_essays'.
3. After step 2 you have to train 3 types of models as discussed below.
4. For all the model use 'auc' (https://scikit-learn.org/stable/modules/model_evaluation.html#roc-metrics) as a metric. check [this \(https://datascience.stackexchange.com/a/20192\)](https://datascience.stackexchange.com/a/20192) for using auc as a metric
5. You are free to choose any number of layers/hidden units but you have to use same type of architectures shown below.
6. You can use any one of the optimizers and choice of Learning rate and momentum, resources: [cs231n class notes \(http://cs231n.github.io/neural-networks-3/\)](http://cs231n.github.io/neural-networks-3/), [cs231n class video \(https://www.youtube.com/watch?v=hd_KFJ5ktUc\)](https://www.youtube.com/watch?v=hd_KFJ5ktUc).
7. For all the model's use [TensorBoard \(https://www.youtube.com/watch?v=2U6Jl7oqRkM\)](https://www.youtube.com/watch?v=2U6Jl7oqRkM) and plot the Metric value and Loss with epoch. While submitting, take a screenshot of plots and include those images in .ipynb notebook and PDF.
8. Use Categorical Cross Entropy as Loss to minimize.

Model-1

Build and Train deep neural network as shown below



ref: <https://i.imgur.com/w395Yk9.png> (<https://i.imgur.com/w395Yk9.png>)

- **Input_seq_total_text_data** --- You have to give Total text data columns. After this use the Embedding layer to get word vectors. Use given predefined glove word vectors, don't train any word vectors. After this use LSTM and get the LSTM output and Flatten that output.
- **Input_school_state** --- Give 'school_state' column as input to embedding layer and Train the Keras Embedding layer.
- **Project_grade_category** --- Give 'project_grade_category' column as input to embedding layer and Train the Keras Embedding layer.
- **Input_clean_categories** --- Give 'input_clean_categories' column as input to embedding layer and Train the Keras Embedding layer.
- **Input_clean_subcategories** --- Give 'input_clean_subcategories' column as input to embedding layer and Train the Keras Embedding layer.
- **Input_clean_subcategories** --- Give 'input_teacher_prefix' column as input to embedding layer and Train the Keras Embedding layer.
- **Input_remaining_teacher_number_of_previously_posted_projects._resource_summary_contains_** --- concatenate remaining columns and add a Dense layer after that.

- For LSTM, you can choose your sequence padding methods on your own or you can train your LSTM without padding, there is no restriction on that.

Below is an example of embedding layer for a categorical columns. In below code all are dummy values, we gave only for reference.

In [2]:

```
# https://stats.stackexchange.com/questions/270546/how-does-keras-embedding-layer-work
input_layer = Input(shape=(n,))
embedding = Embedding(no_1, no_2, input_length=n)(input_layer)
flatten = Flatten()(embedding)
```

```
-----
-
NameError                                Traceback (most recent call last)
<ipython-input-2-ed7dba31d057> in <module>()
      1 # https://stats.stackexchange.com/questions/270546/how-does-keras-embedding-layer-work
----> 2 input_layer = Input(shape=(n,))
      3 embedding = Embedding(no_1, no_2, input_length=n)(input_layer)
      4 flatten = Flatten()(embedding)
```

NameError: name 'Input' is not defined

In [3]:

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

1. Go through this blog, if you have any doubt on using predefined Embedding values in Embedding layer - <https://machinelearningmastery.com/use-word-embedding-layers-deep-learning-keras/> (<https://machinelearningmastery.com/use-word-embedding-layers-deep-learning-keras/>)

2. Please go through this link <https://keras.io/getting-started/functional-api-guide/> (<https://keras.io/getting-started/functional-api-guide/>) and check the 'Multi-input and multi-output models' then you will get to know how to give multiple inputs.

In [4]:

```
import numpy as np
import pandas as pd
from keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential
from keras.layers import Dense, Input, Dropout
from keras.layers import Flatten
from keras.layers import concatenate
from keras.layers.embeddings import Embedding
from keras.models import Model
from keras.utils import to_categorical
from sklearn.model_selection import train_test_split
from keras.preprocessing.text import Tokenizer
import matplotlib.pyplot as plt
import pickle
from keras.layers import LSTM
from keras.preprocessing.text import text_to_word_sequence
import tensorflow as tf
from keras.callbacks import ModelCheckpoint, TensorBoard, ReduceLROnPlateau, EarlyStopping
from keras.layers.normalization import BatchNormalization
from sklearn.feature_extraction.text import TfidfVectorizer
import seaborn as sns
from sklearn.metrics import roc_auc_score
from keras.models import load_model
import tensorflow as tf
```

In [5]:

```
import pandas as pd

data = pd.read_csv("/content/drive/My Drive/9_Donors_choose_DT/preprocessed_data.csv")
data.head()
```

Out[5]:

	school_state	teacher_prefix	project_grade_category	teacher_number_of_previously_posted_projects
0	ca	mrs	grades_prek_2	1
1	ut	ms	grades_3_5	1
2	ca	mrs	grades_prek_2	1
3	ga	mrs	grades_prek_2	1
4	wa	mrs	grades_3_5	1

In [6]:

```
data.columns
```

Out[6]:

```
Index(['school_state', 'teacher_prefix', 'project_grade_category',
      'teacher_number_of_previously_posted_projects', 'project_is_approved',
      'clean_categories', 'clean_subcategories', 'essay', 'price'],
      dtype='object')
```

In [7]:

```
y=data["project_is_approved"].values
```

In [8]:

```
data.drop("project_is_approved",axis = 1, inplace = True)
```

For clean Subcategories

In [9]:

```
d=[]
data['temp']=data['clean_subcategories'].apply(lambda x: x.split()).apply(lambda x: d+x
)
```

In [10]:

```
data['temp']
```

Out[10]:

```
0      [appliedsciences, health_lifescience]
1      [specialneeds]
2      [literacy]
3      [earlydevelopment]
4      [literacy]
...
109243      [teamsports]
109244      [earlydevelopment, other]
109245      [appliedsciences, environmentalscience]
109246      [health_lifescience]
109247      [literacy, literature_writing]
Name: temp, Length: 109248, dtype: object
```

In [11]:

```
d=[]
data['clean_subcategories'].apply(lambda x: x.split()).apply(lambda x: d.extend(x))
z=list(set(d))
from sklearn.preprocessing import LabelEncoder
label = LabelEncoder()
label.fit(z)
```

Out[11]:

```
LabelEncoder()
```

In [12]:

```
data['clean_subcategories_label']=data['temp'].apply(lambda q:label.transform(q))
```

In [13]:

```
data['clean_subcategories_label']
```

Out[13]:

```
0      [0, 14]
1      [26]
2      [17]
3      [6]
4      [17]
...
109243  [27]
109244  [6, 22]
109245  [0, 8]
109246  [14]
109247  [17, 18]
Name: clean_subcategories_label, Length: 109248, dtype: object
```

In [14]:

```
np.max(data['clean_subcategories_label'].apply(lambda x: len(x)))
```

Out[14]:

3

In [15]:

```
from keras.preprocessing import sequence
max_review_length = 3
clean_subcategories_label_ = sequence.pad_sequences(data['clean_subcategories_label'].values, maxlen=max_review_length, padding='post')
```

In [16]:

```
import scipy.sparse as sparse
arr = sparse.coo_matrix(clean_subcategories_label_, shape=(109248,3))
data['clean_subcategories_label_'] = arr.toarray().tolist()
print(data)
```

```
school_state  ... clean_subcategories_label_
0            ca  ...      [0, 14, 0]
1            ut  ...      [26, 0, 0]
2            ca  ...      [17, 0, 0]
3            ga  ...      [6, 0, 0]
4            wa  ...      [17, 0, 0]
...          ...  ...      ...
109243        hi  ...      [27, 0, 0]
109244        nm  ...      [6, 22, 0]
109245        il  ...      [0, 8, 0]
109246        hi  ...      [14, 0, 0]
109247        ca  ...      [17, 18, 0]
```

[109248 rows x 11 columns]

In [17]:

```
data.shape
```

Out[17]:

```
(109248, 11)
```

In [18]:

```
data.drop(['temp', 'clean_subcategories_label'], axis=1, inplace=True)
```

In [19]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(data, y, test_size=0.33, stratify=y)
```

```
print(X_train.shape, Y_train.shape)
print(X_test.shape, Y_test.shape)
```

```
(73196, 9) (73196,)
(36052, 9) (36052,)
```

Tokenizer

In [20]:

```
from keras.preprocessing.text import Tokenizer
tokenizer = Tokenizer(num_words= 1000000 )
tokenizer.fit_on_texts(X_train["essay"])
```

In [21]:

```
X_train['es_tok'] = tokenizer.texts_to_sequences(X_train['essay'].values)
X_test['es_tok'] = tokenizer.texts_to_sequences(X_test['essay'].values)
```

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
"""Entry point for launching an IPython kernel.
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

This is separate from the ipykernel package so we can avoid doing imports until

In [22]:

```
X_test.head()
```

Out[22]:

	school_state	teacher_prefix	project_grade_category	teacher_number_of_previously_p
60249	ga	ms	grades_prek_2	
105342	tn	mrs	grades_prek_2	
43904	ct	mrs	grades_prek_2	
108262	hi	mrs	grades_3_5	
19282	wa	mrs	grades_prek_2	




In [23]:

```
import seaborn as sns
import matplotlib.pyplot as plt

sns.set_theme(style="whitegrid")
plt.hist(X_train['es_tok'].apply(lambda x: len(x)), cumulative=True, density=True, label='CDF', alpha=0.8, color='k', histtype='step')
#np.max(X_train['es_tok'].apply(lambda x: len(x)))
```

Out[23]:

```
(array([0.03113558, 0.37738401, 0.63805399, 0.80541286, 0.90078693,
        0.9558036 , 0.9848489 , 0.99678944, 0.99969944, 1.          ]),
array([ 78. , 104.1, 130.2, 156.3, 182.4, 208.5, 234.6, 260.7, 286.8,
        312.9, 339. ]),
<a list of 1 Patch objects>)
```



Percentile length of review which covers almost 98% is 250, we use Padding of 250

In [24]:

```
from keras.preprocessing import sequence
max_review_length = 250
X_train_pad = sequence.pad_sequences(X_train['es_tok'].values, maxlen=max_review_length)

X_test_pad = sequence.pad_sequences(X_test['es_tok'].values, maxlen=max_review_length)
```

In [25]:

X_train_pad[0]

Out[25]:

```
array([[ 0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0,  0,  0,
        0,  0,  0,  0,  0,  0,  0, 25,  3,
       350, 268,  23, 417,  1, 140,  38,  63, 103,
      1403,  12,  52, 140, 188,  58, 745,  14, 230,
        14, 891, 6035,  1,  33, 1040,  39, 2201, 327,
       177, 2720, 246, 750,  1, 170,  33, 1109,  87,
       462,  53, 230,  14, 171,  8, 122,  40, 412,
       122, 317,  25, 145, 1724, 299, 522, 652, 894,
      1081, 193, 1015, 170, 327,  38, 222,  8, 255,
     1121,  26,  14,  18,  96, 392,  46,  912, 2768,
       312,  71, 774,  1, 1005,  34,  10,  38,  222,
       475,  31, 1276,  5, 395, 527,  1, 12868, 306,
         1, 266,  200, 150,  18, 208,  15,  11, 815,
         1,  15, 882,  67, 684, 403,  71,  74, 3333,
       152, 389,  101, 326, 2408,  55,  13], dtype=int32)
```

In [26]:

```
import pickle
with open('/content/drive/My Drive/9_Donors_choose_DT/glove_vectors', 'rb') as f:
    glove = pickle.load(f)
    glove_words = set(glove.keys())
```

In [27]:

len(tokenizer.word_index.items())

Out[27]:

48129

In [28]:

```
l=tokenizer.word_index.items()
number_of_words_in_corpus = len(tokenizer.word_index)

embedding_matrix = np.zeros((number_of_words_in_corpus+1, 300))
for word, i in l:

    if word in glove_words:
        embedding_matrix[i] =glove[word]
```

In [29]:

```
from sklearn.preprocessing import LabelEncoder
label = LabelEncoder()
label.fit(data['school_state'].values)

X_train_school_state_label_encoding = label.transform(X_train['school_state'].values)
X_test_school_state_label_encoding = label.transform(X_test['school_state'].values)
```

In [30]:

```
data['school_state'].nunique()
```

Out[30]:

51

In [31]:

```
from sklearn.preprocessing import LabelEncoder
label = LabelEncoder()
label.fit(data['teacher_prefix'].values)

X_train_teacher_prefix = label.transform(X_train['teacher_prefix'].values)
X_test_teacher_prefix = label.transform(X_test['teacher_prefix'].values)
```

In [32]:

```
from sklearn.preprocessing import LabelEncoder
label = LabelEncoder()
label.fit(data['clean_categories'].values)

X_train_clean_categories = label.transform(X_train['clean_categories'].values)
X_test_clean_categories = label.transform(X_test['clean_categories'].values)
```

In [33]:

```
from sklearn.preprocessing import LabelEncoder
label = LabelEncoder()
label.fit(data['project_grade_category'].values)

X_train_project_grade_category = label.transform(X_train['project_grade_category'].values)

X_test_project_grade_category = label.transform(X_test['project_grade_category'].values)
```

In [34]:

```
from sklearn.preprocessing import Normalizer
normalizer = Normalizer()
# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.
normalizer.fit(X_train['price'].values.reshape(1,-1))

X_train_price_norm = normalizer.transform(X_train['price'].values.reshape(1,-1))
X_test_price_norm = normalizer.transform(X_test['price'].values.reshape(1,-1))
X_train_price_norm=X_train_price_norm.reshape(-1,1)
X_test_price_norm=X_test_price_norm.reshape(-1,1)
```

In [35]:

```

from sklearn.preprocessing import Normalizer
normalizer = Normalizer()
# normalizer.fit(X_train['price'].values)
# this will rise an error Expected 2D array, got 1D array instead:
# array=[105.22 215.96 96.01 ... 368.98 80.53 709.67].
# Reshape your data either using
# array.reshape(-1, 1) if your data has a single feature
# array.reshape(1, -1) if it contains a single sample.
normalizer.fit(X_train['teacher_number_of_previously_posted_projects'].values.reshape(1, -1))

X_train_teacher_number_of_previously_posted_projects_norm= normalizer.transform(X_train
['teacher_number_of_previously_posted_projects'].values.reshape(1, -1))
X_test_teacher_number_of_previously_posted_projects_norm = normalizer.transform(X_test[
'teacher_number_of_previously_posted_projects'].values.reshape(1, -1))

X_train_teacher_number_of_previously_posted_projects_norm=X_train_teacher_number_of_pre
viously_posted_projects_norm.reshape(-1,1)

X_test_teacher_number_of_previously_posted_projects_norm=X_test_teacher_number_of_previ
ously_posted_projects_norm.reshape(-1,1)
print("After vectorizations")
print(X_train_teacher_number_of_previously_posted_projects_norm.shape, Y_train.shape)
#print(X_cv_teacher_number_of_previously_posted_projects_norm.shape, y_cv.shape)
print(X_test_teacher_number_of_previously_posted_projects_norm.shape, Y_test.shape)
print("=*100)

```

After vectorizations

(73196, 1) (73196,)

(36052, 1) (36052,)

```

=====
=====

```

In [36]:

```

X_train_numeric = np.concatenate((X_train_price_norm , X_train_teacher_number_of_previo
usly_posted_projects_norm) , axis = 1)

X_test_numeric= np.concatenate((X_test_price_norm , X_test_teacher_number_of_previously
_posted_projects_norm) , axis = 1)

```

In [37]:

X_train_numeric.shape

Out[37]:

(73196, 2)

In [38]:

data['clean_subcategories'].nunique()

Out[38]:

401

Reference:<https://stackoverflow.com/questions/57574501/how-to-use-sklearn-auc-in-tensorflow-keras-model-metrics>
(<https://stackoverflow.com/questions/57574501/how-to-use-sklearn-auc-in-tensorflow-keras-model-metrics>)

In [39]:

```
def auc1(y_true, y_pred):  
    if len(np.unique(y_true[:,1])) == 1:  
        return 0.5  
    else:  
        return roc_auc_score(y_true, y_pred)  
  
def auroc(y_true, y_pred):  
    return tf.py_function(auc1, (y_true, y_pred), tf.double)
```

In [40]:

```
from keras.models import Model  
from keras.layers import Input  
from keras.layers import LSTM  
from keras.layers.embeddings import Embedding  
from keras import regularizers  
from keras.regularizers import l2  
from keras.layers import Flatten  
from keras.layers import Dense, Input , Dropout  
from keras.layers import concatenate  
from keras.layers.normalization import BatchNormalization  
from keras.callbacks import TensorBoard
```

In [41]:

```

from keras.models import Model
from keras.layers import Input
from keras.layers import LSTM
from keras.layers.embeddings import Embedding
from keras import regularizers
from keras.regularizers import l2
from keras.layers import Flatten
from keras.layers import Dense, Input , Dropout
from keras.layers import concatenate
from keras.layers.normalization import BatchNormalization
from keras.callbacks import TensorBoard

essay = Input(shape=(250,), name='essay_input')

X = Embedding(output_dim=300, input_dim=number_of_words_in_corpus+1, input_length=250 ,
weights=[embedding_matrix], trainable=False)(essay)
lstm_essay = LSTM(200, recurrent_dropout=0.5, return_sequences=True)(X)
flatten_1 = Flatten()(lstm_essay)

school_state = Input(shape=(1,), name='school_state')
X_school_state = Embedding(output_dim=int(np.sqrt(data['school_state'].nunique())) , in
put_dim=data['school_state'].nunique(), input_length=1)(school_state)
flatten_2 = Flatten()(X_school_state)

teacher_prefix = Input(shape=(1,), name='teacher_prefix')
X_teacher_prefix = Embedding(output_dim=int(np.sqrt(data['teacher_prefix'].nunique()))
, input_dim=data['teacher_prefix'].nunique(), input_length=1)(teacher_prefix)
flatten_3 = Flatten()(X_teacher_prefix)

clean_categories = Input(shape=(1,), name='clean_categories')
X_clean_categories = Embedding(output_dim=int(np.sqrt(data['clean_categories'].nunique
())), input_dim=data['clean_categories'].nunique(), input_length=1)(clean_categories)
flatten_4 = Flatten()(X_clean_categories)

clean_subcategories = Input(shape=(3,), name='clean_subcategories')
X_clean_subcategories = Embedding(output_dim=int(np.sqrt(data['clean_subcategories'].nu
nique())), input_dim=data['clean_subcategories'].nunique(), input_length=3)(clean_subca
tegies)
flatten_5 = Flatten()(X_clean_subcategories)

project_grade_category = Input(shape=(1,), name='project_grade_category')
X_project_grade_category = Embedding(output_dim=int(np.sqrt(data['project_grade_categor
y'].nunique())), input_dim=data['project_grade_category'].nunique(), input_length=1)(pr
oject_grade_category)
flatten_6 = Flatten()(X_project_grade_category)

```

```
numeric_features = Input(shape=(2,) , name="numerical_features")
numeric_dense = Dense(128, activation='relu' , kernel_initializer='he_normal')(numeric_features)

X_concat = concatenate([flatten_1 , flatten_2 , flatten_3 ,flatten_4 , flatten_5 , flatten_6 , numeric_dense])
model = Dense(300, activation="relu", kernel_initializer="he_normal" ,kernel_regularizer=regularizers.l2(0.001))(X_concat)

model = Dropout(0.5)(model)

model = Dense(200,activation="relu",kernel_initializer="glorot_normal")(model)

model = BatchNormalization()(model)

model = Dropout(0.5)(model)

model = Dense(80,activation="relu", kernel_initializer="glorot_normal" )(model)

output = Dense(2, activation='softmax', name='output')(model)

model_1 = Model(inputs=[essay, school_state ,teacher_prefix,clean_categories,
                        clean_subcategories ,project_grade_category ,numeric_features ],
outputs=[output])

print(model_1.summary())
```


WARNING:tensorflow:Layer lstm will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

Model: "functional_1"

Layer (type) to	Output Shape	Param #	Connected
essay_input (InputLayer)	[(None, 250)]	0	
embedding (Embedding) ut[0][0]	(None, 250, 300)	14439000	essay_inp
school_state (InputLayer)	[(None, 1)]	0	
teacher_prefix (InputLayer)	[(None, 1)]	0	
clean_categories (InputLayer)	[(None, 1)]	0	
clean_subcategories (InputLayer)	[(None, 3)]	0	
project_grade_category (InputLa	[(None, 1)]	0	
lstm (LSTM) [0][0]	(None, 250, 200)	400800	embedding
embedding_1 (Embedding) ate[0][0]	(None, 1, 7)	357	school_st
embedding_2 (Embedding) refix[0][0]	(None, 1, 2)	10	teacher_p
embedding_3 (Embedding) egories[0][0]	(None, 1, 7)	357	clean_cat
embedding_4 (Embedding) categories[0][0]	(None, 3, 20)	8020	clean_sub
embedding_5 (Embedding) rade_category[0][0]	(None, 1, 2)	8	project_g
numerical_features (InputLayer)	[(None, 2)]	0	
flatten (Flatten) [0]	(None, 50000)	0	lstm[0]

flatten_1 (Flatten) _1[0][0]	(None, 7)	0	embedding
flatten_2 (Flatten) _2[0][0]	(None, 2)	0	embedding
flatten_3 (Flatten) _3[0][0]	(None, 7)	0	embedding
flatten_4 (Flatten) _4[0][0]	(None, 60)	0	embedding
flatten_5 (Flatten) _5[0][0]	(None, 2)	0	embedding
dense (Dense) _features[0][0]	(None, 128)	384	numerical
concatenate (Concatenate) [0][0]	(None, 50206)	0	flatten
[0][0]			flatten_1
[0][0]			flatten_2
[0][0]			flatten_3
[0][0]			flatten_4
[0][0]			flatten_5
[0]			dense[0]
dense_1 (Dense) te[0][0]	(None, 300)	15062100	concatena
dropout (Dropout) [0][0]	(None, 300)	0	dense_1
dense_2 (Dense) [0][0]	(None, 200)	60200	dropout
batch_normalization (BatchNorma [0][0]	(None, 200)	800	dense_2
dropout_1 (Dropout) malization[0][0]	(None, 200)	0	batch_nor

dense_3 (Dense) [0][0]	(None, 80)	16080	dropout_1
output (Dense) [0][0]	(None, 2)	162	dense_3
=====			
=====			
Total params: 29,988,278			
Trainable params: 15,548,878			
Non-trainable params: 14,439,400			
None			

In [42]:

```
train = [X_train_pad,X_train_school_state_label_encoding.reshape(-1,1),X_train_teacher_prefix.reshape(-1,1),X_train_clean_categories.reshape(-1,1),np.array(X_train['clean_subcategories_label_'].to_list()),X_train_project_grade_category.reshape(-1,1),X_train_numeric]

test = [X_test_pad,X_test_school_state_label_encoding.reshape(-1,1),X_test_teacher_prefix.reshape(-1,1),X_test_clean_categories.reshape(-1,1),np.array(X_test['clean_subcategories_label_'].to_list()),X_test_project_grade_category.reshape(-1,1),X_test_numeric]
```

In [43]:

```
X_train_pad.shape
```

Out[43]:

```
(73196, 250)
```

In [44]:

```
X_train_school_state_label_encoding.reshape(-1,1).shape
```

Out[44]:

```
(73196, 1)
```

In [45]:

```
X_train_teacher_prefix.reshape(-1,1).shape
```

Out[45]:

```
(73196, 1)
```

In [46]:

```
np.array(X_train['clean_subcategories_label_'].to_list()).shape
```

Out[46]:

```
(73196, 3)
```

In [47]:

```
from keras.utils import np_utils

y_train = np_utils.to_categorical(Y_train, 2)
y_test = np_utils.to_categorical(Y_test, 2)
```

In [48]:

y_train

Out[48]:

```
array([[0., 1.],
       [1., 0.],
       [0., 1.],
       ...,
       [0., 1.],
       [0., 1.],
       [0., 1.]], dtype=float32)
```

In [49]:

```
import numpy as np
import tensorflow as tf
from keras.callbacks import Callback
from sklearn.metrics import confusion_matrix, f1_score, precision_score, recall_score, roc_auc_score

class Metrics(tf.keras.callbacks.Callback):
    def __init__(self, validation_data=None, validation_target=None):
        super(Metrics, self).__init__()
        self.validation_data = validation_data
        self.validation_target=validation_target
        # best_weights to store the weights at which the minimum loss occurs.

    def on_train_begin(self, logs={}):
        self.val_f1s=[]
        self.val_auc=[]

    def on_epoch_end(self, epoch, logs={}):

        val_predict =np.array((self.model.predict(self.validation_data)))
        val_targ = np.array(self.validation_target,dtype=int)
        #_val_f1 = f1_score(val_targ, val_predict)
        _val_auc=roc_auc_score(val_targ[:,0], val_predict[:,0],average='macro')
        # self.val_f1s.append(_val_f1)
        #self.val_auc.append(_val_auc)
        print(' -val_auc_score: '+str(_val_auc))
```

In [50]:

```
#Input Layer
import warnings
warnings.filterwarnings("ignore")
import datetime
import os
checkpoint1 = ModelCheckpoint("model_1.h5",

                               monitor="val_loss",
                               mode="auto",
                               save_best_only = True,
                               verbose=1)
earlystop1 = EarlyStopping(monitor = 'val_loss',
                           mode="auto",
                           min_delta = 0,
                           patience = 4,
                           verbose = 2)

log_dir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,histogram_freq=1,
write_graph=True,write_grads=True)

metric=Metrics(test, y_test)

callbacks_1= [checkpoint1,earlystop1,tensorboard_callback]
```

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the
`TensorBoard` Callback.

In [51]:

```
model_1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=[auroc])
history1 = model_1.fit(train, y_train, batch_size=512, epochs=10, verbose=1, callbacks=c
allbacks_1, validation_data=(test, y_test))
```

Epoch 1/10

1/143 [.....] - ETA: 0s - loss: 2.5252 - auroc: 0.4815
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/summary_ops_v2.py:1277: stop (from tensorflow.python.eager.profiler) is deprecated and will be removed after 2020-07-01.

Instructions for updating:

use `tf.profiler.experimental.stop` instead.

2/143 [.....] - ETA: 5:51 - loss: 2.4828 - auroc: 0.4741
WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared to the batch time (batch time: 1.3537s vs `on_train_batch_end` time: 3.6372s). Check your callbacks.

143/143 [=====] - ETA: 0s - loss: 0.9859 - auroc: 0.5491

Epoch 00001: val_loss improved from inf to 0.59027, saving model to model_1.h5

143/143 [=====] - 142s 991ms/step - loss: 0.9859 - auroc: 0.5491 - val_loss: 0.5903 - val_auroc: 0.6661

Epoch 2/10

143/143 [=====] - ETA: 0s - loss: 0.5503 - auroc: 0.6541

Epoch 00002: val_loss improved from 0.59027 to 0.49614, saving model to model_1.h5

143/143 [=====] - 135s 942ms/step - loss: 0.5503 - auroc: 0.6541 - val_loss: 0.4961 - val_auroc: 0.7167

Epoch 3/10

143/143 [=====] - ETA: 0s - loss: 0.4909 - auroc: 0.7004

Epoch 00003: val_loss improved from 0.49614 to 0.47218, saving model to model_1.h5

143/143 [=====] - 131s 917ms/step - loss: 0.4909 - auroc: 0.7004 - val_loss: 0.4722 - val_auroc: 0.7304

Epoch 4/10

143/143 [=====] - ETA: 0s - loss: 0.4617 - auroc: 0.7235

Epoch 00004: val_loss improved from 0.47218 to 0.44821, saving model to model_1.h5

143/143 [=====] - 137s 959ms/step - loss: 0.4617 - auroc: 0.7235 - val_loss: 0.4482 - val_auroc: 0.7408

Epoch 5/10

143/143 [=====] - ETA: 0s - loss: 0.4447 - auroc: 0.7353

Epoch 00005: val_loss improved from 0.44821 to 0.44606, saving model to model_1.h5

143/143 [=====] - 140s 976ms/step - loss: 0.4447 - auroc: 0.7353 - val_loss: 0.4461 - val_auroc: 0.7387

Epoch 6/10

143/143 [=====] - ETA: 0s - loss: 0.4397 - auroc: 0.7462

Epoch 00006: val_loss improved from 0.44606 to 0.44260, saving model to model_1.h5

143/143 [=====] - 142s 992ms/step - loss: 0.4397 - auroc: 0.7462 - val_loss: 0.4426 - val_auroc: 0.7496

Epoch 7/10

143/143 [=====] - ETA: 0s - loss: 0.4410 - auroc: 0.7548

Epoch 00007: val_loss did not improve from 0.44260

143/143 [=====] - 138s 965ms/step - loss: 0.4410 - auroc: 0.7548 - val_loss: 0.4605 - val_auroc: 0.7530

Epoch 8/10

143/143 [=====] - ETA: 0s - loss: 0.4365 - auroc: 0.7592

Epoch 00008: val_loss improved from 0.44260 to 0.44257, saving model to model_1.h5

143/143 [=====] - 140s 981ms/step - loss: 0.4365
- auroc: 0.7592 - val_loss: 0.4426 - val_auroc: 0.7530

Epoch 9/10

143/143 [=====] - ETA: 0s - loss: 0.4375 - auroc: 0.7719

Epoch 00009: val_loss did not improve from 0.44257

143/143 [=====] - 138s 967ms/step - loss: 0.4375
- auroc: 0.7719 - val_loss: 0.4427 - val_auroc: 0.7512

Epoch 10/10

143/143 [=====] - ETA: 0s - loss: 0.4395 - auroc: 0.7748

Epoch 00010: val_loss did not improve from 0.44257

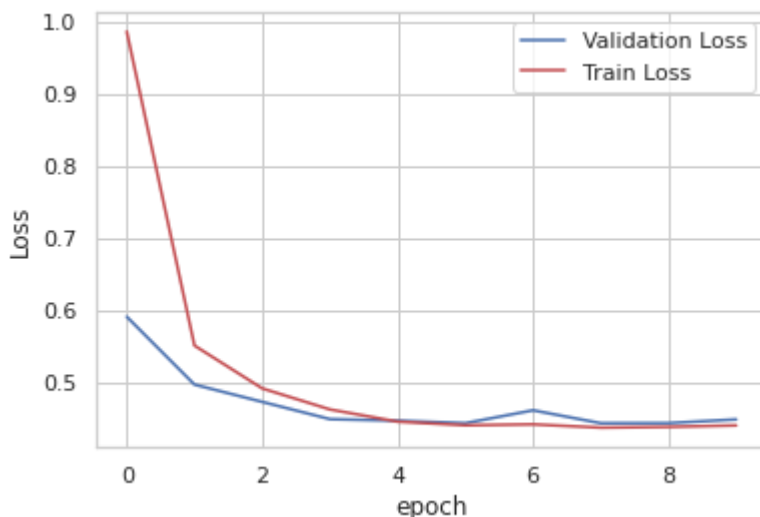
143/143 [=====] - 138s 968ms/step - loss: 0.4395
- auroc: 0.7748 - val_loss: 0.4479 - val_auroc: 0.7531

In [52]:

```
%matplotlib inline
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Loss')
x = list(range(0,10))
vy = history1.history['val_loss']
ty = history1.history['loss']
ax.plot(x, vy, 'b', label="Validation Loss")
ax.plot(x, ty, 'r', label="Train Loss")
plt.legend()
```


Out[52]:

<matplotlib.legend.Legend at 0x7f09f688fbe0>



In [53]:

```
plt.plot(history1.history['auroc'], 'r')
plt.plot(history1.history['val_auroc'], 'b')
plt.legend({'Train ROCAUC': 'r', 'Test ROCAUC': 'b'})
plt.show()
```



In [54]:

```
dot_img_file = '/tmp/model_1.png'
tf.keras.utils.plot_model(model_1, to_file=dot_img_file, show_shapes=True)
```

Out[54]:



Model_2

Model-2

Use the same model as above but for 'input_seq_total_text_data' give only some words in the sentence not all the words. Filter the words as below.

1. Train the TF-IDF on the Train data
2. Get the idf value for each word we have in the train data.
3. Remove the low idf value and high idf value words from our data. Do some analysis on the Idf values and based on those values choose the low and high threshold value. Because very frequent words and very very rare words don't give much information. (you can plot a box plots and take only the idf scores within IQR range and corresponding words)
4. Train the LSTM after removing the Low and High idf value words. (In model-1 Train on total data but in Model-2 train on data after removing some words based on IDF values)

In [55]:

```
vectorizer = TfidfVectorizer()  
vectorizer.fit_transform(X_train["essay"])
```

Out[55]:

```
<73196x48093 sparse matrix of type '<class 'numpy.float64'>'  
  with 7908625 stored elements in Compressed Sparse Row format>
```

Depends upon idf percentile value, we are going to take specific words

In [56]:

```

arr=vectorizer.idf_
print("10th percentile of arr : ",
      np.percentile(arr, 10))
print("40th percentile of arr : ",
      np.percentile(arr, 40))
print("50th percentile of arr : ",
      np.percentile(arr, 50))
print("25th percentile of arr : ",
      np.percentile(arr, 25))
print("75th percentile of arr : ",
      np.percentile(arr, 75))
print("90th percentile of arr : ",
      np.percentile(arr, 90))
print("95th percentile of arr : ",
      np.percentile(arr, 95))
print("98th percentile of arr : ",
      np.percentile(arr, 98))
print("99th percentile of arr : ",
      np.percentile(arr, 99))

```

```

10th percentile of arr : 7.473521896790654
40th percentile of arr : 10.591471803068893
50th percentile of arr : 11.102297426834886
25th percentile of arr : 9.492859514400784
75th percentile of arr : 11.50776253494305
90th percentile of arr : 11.507762534943051
95th percentile of arr : 11.50776253494305
98th percentile of arr : 11.50776253494305
99th percentile of arr : 11.50776253494305

```

In [57]:

```

mid_idf=[str(i) for i,j in zip(vectorizer.get_feature_names() ,vectorizer.idf_) if j>2
and j<11.5]

```

In [58]:

```
len(set(mid_idf))
```

Out[58]:

29389

In [59]:

```

tokenizer = Tokenizer(num_words= 100000 )
tokenizer.fit_on_texts(mid_idf)

X_train['essay_tok_mid'] = tokenizer.texts_to_sequences(X_train['essay'].values)
X_test['essay_tok_mid'] = tokenizer.texts_to_sequences(X_test['essay'].values)

```


In [60]:

```
import seaborn as sns
import matplotlib.pyplot as plt

sns.set_theme(style="whitegrid")
plt.hist(X_train['essay_tok_mid'].apply(lambda x: len(x)), cumulative=True, density=True,
label='CDF', alpha=0.8, color='k', histtype='step')
#np.max(X_train['es_tok'].apply(lambda x: len(x)))
```

Out[60]:

```
(array([0.00885294, 0.29154599, 0.62286191, 0.8147713 , 0.91685338,
        0.96917864, 0.99147494, 0.99866113, 0.99986338, 1.          ]),
 array([ 43. ,  68.1,  93.2, 118.3, 143.4, 168.5, 193.6, 218.7, 243.8,
        268.9, 294. ]),
 <a list of 1 Patch objects>)
```



In [101]:

```
from keras.preprocessing import sequence
max_review_length = 200
X_train_pad1 = sequence.pad_sequences(X_train['essay_tok_mid'].values, maxlen=max_review_length)

X_test_pad1 = sequence.pad_sequences(X_test['essay_tok_mid'].values, maxlen=max_review_length)
```

In [62]:

```
X_train_pad1.shape
```

Out[62]:

```
(73196, 200)
```

In [117]:

```

from keras.models import Model
from keras.layers import Input
from keras.layers import LSTM
from keras.layers.embeddings import Embedding
from keras import regularizers
from keras.regularizers import l2
from keras.layers import Flatten
from keras.layers import Dense, Input , Dropout
from keras.layers import concatenate
from keras.layers.normalization import BatchNormalization
from keras.callbacks import TensorBoard

essay1 = Input(shape=(200,), name='essay_input1')

X = Embedding(len(mid_idf)+1,300,input_length=200,trainable=True)(essay1)
lstm_essay = LSTM(100,recurrent_dropout=0.5,return_sequences=True,kernel_regularizer=regularizers.l2(0.001))(X)
flatten_1_mid = Flatten()(lstm_essay)

school_state1 = Input(shape=(1,), name='school_state1')
X_school_state = Embedding(output_dim=int(np.sqrt(data['school_state'].nunique())) , input_dim=data['school_state'].nunique(), input_length=1)(school_state1)
flatten_2 = Flatten()(X_school_state)

teacher_prefix1 = Input(shape=(1,), name='teacher_prefix1')
X_teacher_prefix = Embedding(output_dim=int(np.sqrt(data['teacher_prefix'].nunique())) , input_dim=data['teacher_prefix'].nunique(), input_length=1)(teacher_prefix1)
flatten_3 = Flatten()(X_teacher_prefix)

clean_categories1 = Input(shape=(1,), name='clean_categories1')
X_clean_categories = Embedding(output_dim=int(np.sqrt(data['clean_categories'].nunique())) , input_dim=data['clean_categories'].nunique(), input_length=1)(clean_categories1)
flatten_4 = Flatten()(X_clean_categories)

clean_subcategories1 = Input(shape=(3,), name='clean_subcategories1')
X_clean_subcategories = Embedding(output_dim=int(np.sqrt(data['clean_subcategories'].nunique())) , input_dim=data['clean_subcategories'].nunique(), input_length=3)(clean_subcategories1)
flatten_5 = Flatten()(X_clean_subcategories)

project_grade_category1 = Input(shape=(1,), name='project_grade_category1')
X_project_grade_category = Embedding(output_dim=int(np.sqrt(data['project_grade_category'].nunique())) , input_dim=data['project_grade_category'].nunique(), input_length=1)(project_grade_category1)
flatten_6 = Flatten()(X_project_grade_category)

```

```
numeric_features1 = Input(shape=(2,) , name="numerical_features1")
numeric_dense = Dense(128, activation='relu' , kernel_initializer='he_normal')(numeric_
features1)

X_concat = concatenate([flatten_1_mid , flatten_2 , flatten_3 ,flatten_4 , flatten_5 ,
flatten_6 , numeric_dense])
model = Dense(50, activation="relu", kernel_initializer="he_normal")(X_concat)

model = Dropout(0.5)(model)

model = Dense(100,activation="relu",kernel_initializer="glorot_normal")(model)

model = BatchNormalization()(model)

model = Dropout(0.5)(model)

model = Dense(80,activation="sigmoid", kernel_initializer="glorot_normal")(model)

output1 = Dense(2, activation='softmax', name='output1')(model)

model_2 = Model(inputs=[essay1, school_state1 ,teacher_prefix1,clean_categories1,
                        clean_subcategories1 ,project_grade_category1 ,numeric_features1
],outputs=[output1])

print(model_2.summary())
```

WARNING:tensorflow:Layer lstm_12 will not use cuDNN kernel since it does n't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

Model: "functional_25"

Layer (type)	Output Shape	Param #	Connected to
essay_input1 (InputLayer)	[(None, 200)]	0	
embedding_57 (Embedding) input1[0][0]	(None, 200, 300)	8817000	essay_i
school_state1 (InputLayer)	[(None, 1)]	0	
teacher_prefix1 (InputLayer)	[(None, 1)]	0	
clean_categories1 (InputLayer)	[(None, 1)]	0	
clean_subcategories1 (InputLayer)	[(None, 3)]	0	
project_grade_category1 (InputLayer)	[(None, 1)]	0	
lstm_12 (LSTM) ng_57[0][0]	(None, 200, 100)	160400	embeddi
embedding_58 (Embedding) state1[0][0]	(None, 1, 7)	357	school_
embedding_59 (Embedding) _prefix1[0][0]	(None, 1, 2)	10	teacher
embedding_60 (Embedding) ategories1[0][0]	(None, 1, 7)	357	clean_c
embedding_61 (Embedding) ubcategories1[0][0]	(None, 3, 20)	8020	clean_s
embedding_62 (Embedding) _grade_category1[0][0]	(None, 1, 2)	8	project
numerical_features1 (InputLayer)	[(None, 2)]	0	
flatten_60 (Flatten) [0][0]	(None, 20000)	0	lstm_12

flatten_61 (Flatten) ng_58[0][0]	(None, 7)	0	embeddi
flatten_62 (Flatten) ng_59[0][0]	(None, 2)	0	embeddi
flatten_63 (Flatten) ng_60[0][0]	(None, 7)	0	embeddi
flatten_64 (Flatten) ng_61[0][0]	(None, 60)	0	embeddi
flatten_65 (Flatten) ng_62[0][0]	(None, 2)	0	embeddi
dense_45 (Dense) al_features1[0][0]	(None, 128)	384	numeric
concatenate_12 (Concatenate) _60[0][0]	(None, 20206)	0	flatten
_61[0][0]			flatten
_62[0][0]			flatten
_63[0][0]			flatten
_64[0][0]			flatten
_65[0][0]			flatten
5[0][0]			dense_4
dense_46 (Dense) nate_12[0][0]	(None, 50)	1010350	concate
dropout_24 (Dropout) 6[0][0]	(None, 50)	0	dense_4
dense_47 (Dense) _24[0][0]	(None, 100)	5100	dropout
batch_normalization_12 (BatchNo 7[0][0]	(None, 100)	400	dense_4
dropout_25 (Dropout) ormalization_12[0][0]	(None, 100)	0	batch_n

dense_48 (Dense) _25[0][0]	(None, 80)	8080	dropout
output1 (Dense) 8[0][0]	(None, 2)	162	dense_4

=====

Total params: 10,010,628
Trainable params: 10,010,428
Non-trainable params: 200

None

In [103]:

```
train = [X_train_pad1,X_train_school_state_label_encoding.reshape(-1,1),X_train_teacher_prefix.reshape(-1,1),X_train_clean_categories.reshape(-1,1),np.array(X_train['clean_subcategories_label_'].to_list()),X_train_project_grade_category.reshape(-1,1),X_train_numeric]

test = [X_test_pad1,X_test_school_state_label_encoding.reshape(-1,1),X_test_teacher_prefix.reshape(-1,1),X_test_clean_categories.reshape(-1,1),np.array(X_test['clean_subcategories_label_'].to_list()),X_test_project_grade_category.reshape(-1,1),X_test_numeric]
```

In [104]:

```
from keras.utils import np_utils

y_train = np_utils.to_categorical(Y_train, 2)
y_test = np_utils.to_categorical(Y_test, 2)
```

In [66]:

```
y_test.shape
```

Out[66]:

(36052, 2)

In [66]:

In [118]:

```
#Input Layer
import warnings
warnings.filterwarnings("ignore")
import datetime
import os
checkpoint1 = ModelCheckpoint("model_2.h5",

                               monitor="val_loss",
                               mode="auto",
                               save_best_only = True,
                               verbose=1)
earlystop1 = EarlyStopping(monitor = 'val_loss',
                           mode="auto",
                           min_delta = 0,
                           patience = 2,
                           verbose = 2)

log_dir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,histogram_freq=1,
write_graph=True,write_grads=True)
metric=Metrics(test, y_test)

callbacks_2= [checkpoint1,earlystop1,tensorboard_callback]
```

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the
`TensorBoard` Callback.

In [119]:

```
model_2.compile(optimizer='Nadam', loss='categorical_crossentropy', metrics=[auroc])
history = model_2.fit(train, y_train, batch_size=600, epochs=10, verbose=1, callbacks=callbacks_2, validation_data=(test, y_test))
```

Epoch 1/10

2/122 [.....] - ETA: 4:43 - loss: 0.7612 - auroc: 0.4918WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared to the batch time (batch time: 1.3319s vs `on_train_batch_end` time: 3.3977s). Check your callbacks.

122/122 [=====] - ETA: 0s - loss: 0.5081 - auroc: 0.6630

Epoch 00001: val_loss improved from inf to 0.41136, saving model to model_2.h5

122/122 [=====] - 120s 982ms/step - loss: 0.5081 - auroc: 0.6630 - val_loss: 0.4114 - val_auroc: 0.7315

Epoch 2/10

122/122 [=====] - ETA: 0s - loss: 0.3727 - auroc: 0.7649

Epoch 00002: val_loss improved from 0.41136 to 0.38823, saving model to model_2.h5

122/122 [=====] - 116s 954ms/step - loss: 0.3727 - auroc: 0.7649 - val_loss: 0.3882 - val_auroc: 0.7314

Epoch 3/10

122/122 [=====] - ETA: 0s - loss: 0.3467 - auroc: 0.7978

Epoch 00003: val_loss did not improve from 0.38823

122/122 [=====] - 115s 945ms/step - loss: 0.3467 - auroc: 0.7978 - val_loss: 0.3943 - val_auroc: 0.7268

Epoch 4/10

122/122 [=====] - ETA: 0s - loss: 0.3231 - auroc: 0.8293

Epoch 00004: val_loss did not improve from 0.38823

122/122 [=====] - 114s 938ms/step - loss: 0.3231 - auroc: 0.8293 - val_loss: 0.4114 - val_auroc: 0.7101


Epoch 00004: early stopping

In [120]:

```
%matplotlib inline
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Loss')
x = list(range(0,4))
vy = history.history['val_loss']
ty = history.history['loss']
ax.plot(x, vy, 'b', label="Validation Loss")
ax.plot(x, ty, 'r', label="Train Loss")
plt.legend()
```

Out[120]:

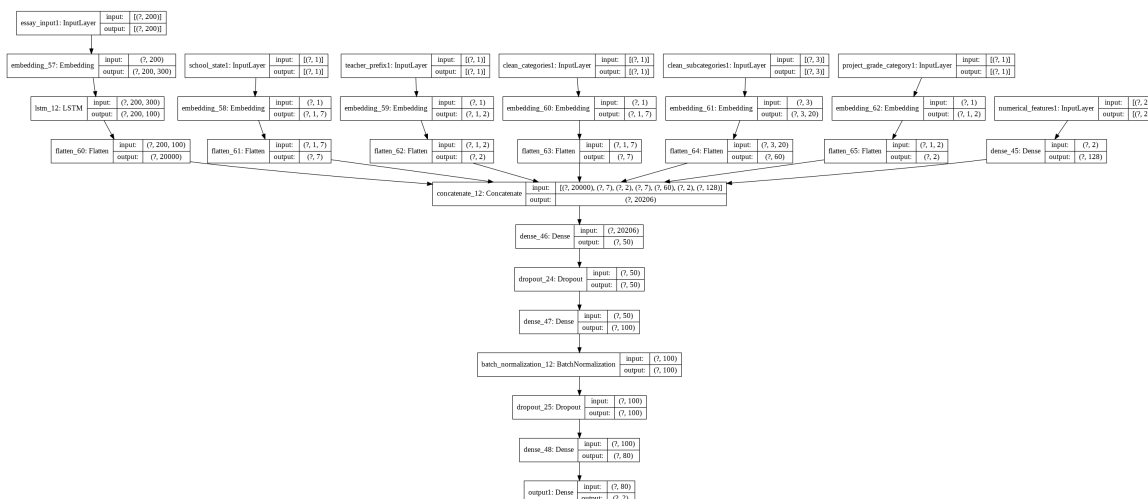
<matplotlib.legend.Legend at 0x7f0859e87390>



In [121]:


```
dot_img_file = '/tmp/model_2.png'
tf.keras.utils.plot_model(model_2, to_file=dot_img_file, show_shapes=True)
```

Out[121]:



Model_3

Model-3



ref: <https://i.imgur.com/fkQ8nGo.png> (<https://i.imgur.com/fkQ8nGo.png>)

- **input_seq_total_text_data:**

- . Use text column('essay'), and use the Embedding layer to get word vectors.
- . Use given predefined glove word vectors, don't train any word vectors.
- . Use LSTM that is given above, get the LSTM output and Flatten that output.
- . You are free to preprocess the input text as you needed.

- **Other_than_text_data:**

- . Convert all your Categorical values to onehot coded and then concatenate all these onehot vectors
- . Neumerical values and use [CNN1D \(https://keras.io/getting-started/sequential-model-guide/#sequence-classification-with-1d-convolutions\)](https://keras.io/getting-started/sequential-model-guide/#sequence-classification-with-1d-convolutions) as shown in above figure.
- . You are free to choose all CNN parameters like kernel sizes, stride.

</pre>

In [73]:

```
from keras.layers import Conv1D
from sklearn.feature_extraction.text import CountVectorizer
```

In [74]:

```
vectorizer = CountVectorizer()
vectorizer.fit(X_train['school_state'].values)
x_train_state_one_hot = vectorizer.transform(X_train['school_state'].values)
x_test_state_one_hot = vectorizer.transform(X_test['school_state'].values)
```

```
print(x_train_state_one_hot.shape, y_train.shape)
print(x_test_state_one_hot.shape, y_test.shape)
```

```
(73196, 51) (73196, 2)
(36052, 51) (36052, 2)
```


In [75]:

```
vectorizer = CountVectorizer()
vectorizer.fit(X_train['clean_categories'].values)

x_train_categories_one_hot = vectorizer.transform(X_train['clean_categories'].values)
x_test_categories_one_hot = vectorizer.transform(X_test['clean_categories'].values)


print(x_train_categories_one_hot.shape, y_train.shape)
print(x_test_categories_one_hot.shape, y_test.shape)
```

```
(73196, 9) (73196, 2)
(36052, 9) (36052, 2)
```

In [76]:

```
vectorizer = CountVectorizer()
vectorizer.fit(X_train['clean_subcategories'].values)

x_train_subcategories_one_hot = vectorizer.transform(X_train['clean_subcategories'].values)
x_test_subcategories_one_hot = vectorizer.transform(X_test['clean_subcategories'].values)


print(x_train_subcategories_one_hot.shape, y_train.shape)
print(x_test_subcategories_one_hot.shape, y_test.shape)
```

```
(73196, 30) (73196, 2)
(36052, 30) (36052, 2)
```

In [77]:

```
vectorizer = CountVectorizer()
vectorizer.fit(X_train['teacher_prefix'].values)
x_train_teacher_prefix_one_hot = vectorizer.transform(X_train['teacher_prefix'].values)
x_test_teacher_prefix_one_hot = vectorizer.transform(X_test['teacher_prefix'].values)


print(x_train_teacher_prefix_one_hot.shape, y_train.shape)
print(x_test_teacher_prefix_one_hot.shape, y_test.shape)
```

```
(73196, 5) (73196, 2)
(36052, 5) (36052, 2)
```

In [78]:

```
vectorizer = CountVectorizer()
vectorizer.fit(X_train['project_grade_category'].values)
x_train_project_grade_one_hot = vectorizer.transform(X_train['project_grade_category'].values)
x_test_project_grade_one_hot = vectorizer.transform(X_test['project_grade_category'].values)

print(x_train_project_grade_one_hot.shape, y_train.shape)
print(x_test_project_grade_one_hot.shape, y_test.shape)

(73196, 4) (73196, 2)
(36052, 4) (36052, 2)
```

In [79]:

```
from scipy.sparse import hstack
```

In [80]:

```
train_features_wot= hstack((x_train_project_grade_one_hot,x_train_teacher_prefix_one_hot,x_train_categories_one_hot,x_train_subcategories_one_hot,x_train_state_one_hot,X_train_price_norm,X_train_teacher_number_of_previously_posted_projects_norm)).todense()
test_features_wot = hstack((x_test_project_grade_one_hot,x_test_teacher_prefix_one_hot,x_test_categories_one_hot,x_test_subcategories_one_hot,x_test_state_one_hot,X_test_price_norm,X_test_teacher_number_of_previously_posted_projects_norm)).todense()
```

In [81]:

```
rest_train = np.expand_dims(train_features_wot,2)
rest_test = np.expand_dims(test_features_wot,2)
```

In [82]:

```
print("train data shape",rest_train.shape)
print("test data shape",rest_test.shape)
```

```
train data shape (73196, 101, 1)
test data shape (36052, 101, 1)
```

In [83]:

```
from keras.preprocessing.text import Tokenizer
tokenizer = Tokenizer(num_words= 1000000 )
tokenizer.fit_on_texts(X_train["essay"])
```

In [84]:

```
X_train['es_tok'] = tokenizer.texts_to_sequences(X_train['essay'].values)
X_test['es_tok'] = tokenizer.texts_to_sequences(X_test['essay'].values)
```

In [85]:

```
l=tokenizer.word_index.items()
number_of_words_in_corpus = len(tokenizer.word_index)

embedding_matrix = np.zeros((number_of_words_in_corpus+1, 300))
for word, i in l:

    if word in glove_words:
        embedding_matrix[i] =glove[word]
```

In [86]:

```
from keras.preprocessing import sequence
max_review_length = 250
X_train_pad = sequence.pad_sequences(X_train['es_tok'].values, maxlen=max_review_length
)

X_test_pad  = sequence.pad_sequences(X_test['es_tok'].values, maxlen=max_review_length)
```

In [86]:

In [93]:

```

from keras.layers import Conv1D
from keras.initializers import he_normal

essay = Input(shape=(250,))

X = Embedding(output_dim=300, input_dim=number_of_words_in_corpus+1, input_length=300 ,
weights=[embedding_matrix],trainable=False)(essay)
lstm_essay = LSTM(200,recurrent_dropout=0.5,return_sequences=True)(X)
flatten_1 = Flatten()(lstm_essay)


input_wot = Input(shape=(101,1))

con = Conv1D(300 , 3 , activation='relu' , kernel_initializer=he_normal(seed=10) , padding='valid')(input_wot)

convo = Conv1D(150 , 3 , activation='relu' , kernel_initializer=he_normal(seed=0) , padding='valid')(con)

flatten_2 = Flatten()(convo)


x_concat = concatenate([flatten_1 , flatten_2])

x = Dense(120, activation="relu", kernel_initializer="he_normal",kernel_regularizer=regularizers.l2(0.001) )(x_concat)

x=Dropout(0.5)(x)

x = Dense(200,activation="sigmoid",kernel_initializer="glorot_normal",kernel_regularizer=regularizers.l2(0.001) )(x)

x = BatchNormalization()(x)

x=Dropout(0.5)(x)

x = Dense(75,activation="relu", kernel_initializer="he_normal",kernel_regularizer=regularizers.l2(0.001))(x)


output = Dense(2, activation='softmax', name='output')(x)

model_3= Model(inputs=[essay, input_wot],outputs=[output])

print(model_3.summary())

```

WARNING:tensorflow:Layer lstm_6 will not use cuDNN kernel since it does n't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

Model: "functional_13"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_5 (InputLayer)	[(None, 250)]	0	
=====			
input_6 (InputLayer)	[(None, 101, 1)]	0	
=====			
embedding_26 (Embedding) [0][0]	(None, 250, 300)	14439000	input_5
=====			
conv1d_4 (Conv1D) [0][0]	(None, 99, 300)	1200	input_6
=====			
lstm_6 (LSTM) ng_26[0][0]	(None, 250, 200)	400800	embedding_26[0][0]
=====			
conv1d_5 (Conv1D) 4[0][0]	(None, 97, 150)	135150	conv1d_4[0][0]
=====			
flatten_28 (Flatten) [0][0]	(None, 50000)	0	lstm_6
=====			
flatten_29 (Flatten) 5[0][0]	(None, 14550)	0	conv1d_5[0][0]
=====			
concatenate_6 (Concatenate) _28[0][0]	(None, 64550)	0	flatten_28[0][0]
=====			
dense_22 (Dense) nate_6[0][0]	(None, 120)	7746120	concatenate_6[0][0]
=====			
dropout_12 (Dropout) 2[0][0]	(None, 120)	0	dense_22[0][0]
=====			
dense_23 (Dense) _12[0][0]	(None, 200)	24200	dropout_12[0][0]
=====			
batch_normalization_6 (BatchNormaliz 3[0][0]	(None, 200)	800	dense_23[0][0]

dropout_13 (Dropout) ormalization_6[0][0]	(None, 200)	0	batch_n
dense_24 (Dense) _13[0][0]	(None, 75)	15075	dropout
output (Dense) 4[0][0]	(None, 2)	152	dense_2

=====

Total params: 22,762,497
Trainable params: 8,323,097
Non-trainable params: 14,439,400

None

In [88]:

```
train = [X_train_pad,rest_train]
test = [X_test_pad,rest_test]
```

In [89]:

```
#Input Layer
import warnings
warnings.filterwarnings("ignore")
import datetime
import os
checkpoint1 = ModelCheckpoint("model_3.h5",

                               monitor="val_loss",
                               mode="auto",
                               save_best_only = True,
                               verbose=1)
earlystop1 = EarlyStopping(monitor = 'val_loss',
                            mode="auto",
                            min_delta = 0,
                            patience = 4,
                            verbose = 2)

log_dir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,histogram_freq=1,
write_graph=True,write_grads=True)
metric=Metrics(test, y_test)

callbacks_2= [checkpoint1,earlystop1,tensorboard_callback]
```

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the
`TensorBoard` Callback.

In [95]:

```
model_3.compile(optimizer='Nadam', loss='categorical_crossentropy', metrics=[auroc])  
history2= model_3.fit(train, y_train, batch_size=600, epochs=20, verbose=2, callbacks=ca  
llbacks_2, validation_data=(test, y_test))
```

Epoch 1/20

WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared to the batch time (batch time: 1.3865s vs `on_train_batch_end` time: 3.8618s). Check your callbacks.

Epoch 00001: val_loss did not improve from 0.42303

122/122 - 130s - loss: 0.7631 - auroc: 0.6179 - val_loss: 0.7736 - val_auroc: 0.6781

Epoch 2/20

Epoch 00002: val_loss did not improve from 0.42303

122/122 - 125s - loss: 0.5909 - auroc: 0.7054 - val_loss: 0.6728 - val_auroc: 0.7203

Epoch 3/20

Epoch 00003: val_loss did not improve from 0.42303

122/122 - 125s - loss: 0.5236 - auroc: 0.7301 - val_loss: 0.6177 - val_auroc: 0.7386

Epoch 4/20

Epoch 00004: val_loss did not improve from 0.42303

122/122 - 125s - loss: 0.4835 - auroc: 0.7497 - val_loss: 0.5959 - val_auroc: 0.7504

Epoch 5/20

Epoch 00005: val_loss did not improve from 0.42303

122/122 - 125s - loss: 0.4563 - auroc: 0.7615 - val_loss: 0.5511 - val_auroc: 0.7518

Epoch 6/20

Epoch 00006: val_loss did not improve from 0.42303

122/122 - 125s - loss: 0.4412 - auroc: 0.7697 - val_loss: 0.5459 - val_auroc: 0.7506

Epoch 7/20

Epoch 00007: val_loss did not improve from 0.42303

122/122 - 126s - loss: 0.4291 - auroc: 0.7849 - val_loss: 0.4818 - val_auroc: 0.7580

Epoch 8/20

Epoch 00008: val_loss did not improve from 0.42303

122/122 - 125s - loss: 0.4210 - auroc: 0.7955 - val_loss: 0.5625 - val_auroc: 0.7539

Epoch 9/20

Epoch 00009: val_loss did not improve from 0.42303

122/122 - 124s - loss: 0.4157 - auroc: 0.8121 - val_loss: 0.5212 - val_auroc: 0.7536

Epoch 10/20

Epoch 00010: val_loss did not improve from 0.42303

122/122 - 123s - loss: 0.4724 - auroc: 0.5543 - val_loss: 0.4656 - val_auroc: 0.6632

Epoch 11/20

Epoch 00011: val_loss did not improve from 0.42303

122/122 - 125s - loss: 0.4316 - auroc: 0.6194 - val_loss: 0.5496 - val_auroc: 0.7419

Epoch 12/20

Epoch 00012: val_loss did not improve from 0.42303

122/122 - 125s - loss: 0.3825 - auroc: 0.7955 - val_loss: 0.4828 - val_auroc: 0.7475

Epoch 13/20

Epoch 00013: val_loss did not improve from 0.42303

122/122 - 124s - loss: 0.3762 - auroc: 0.8304 - val_loss: 0.5603 - val_auroc: 0.7483

Epoch 14/20

Epoch 00014: val_loss did not improve from 0.42303

122/122 - 125s - loss: 0.3775 - auroc: 0.8518 - val_loss: 0.4896 - val_auroc: 0.7338


Epoch 00014: early stopping

In [97]:

```
%matplotlib inline
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Loss')
x = list(range(0,14))
vy = history2.history['val_loss']
ty = history2.history['loss']
ax.plot(x, vy, 'b', label="Validation Loss")
ax.plot(x, ty, 'r', label="Train Loss")
plt.legend()
```


Out[97]:

<matplotlib.legend.Legend at 0x7f0716f7a630>



In [99]:

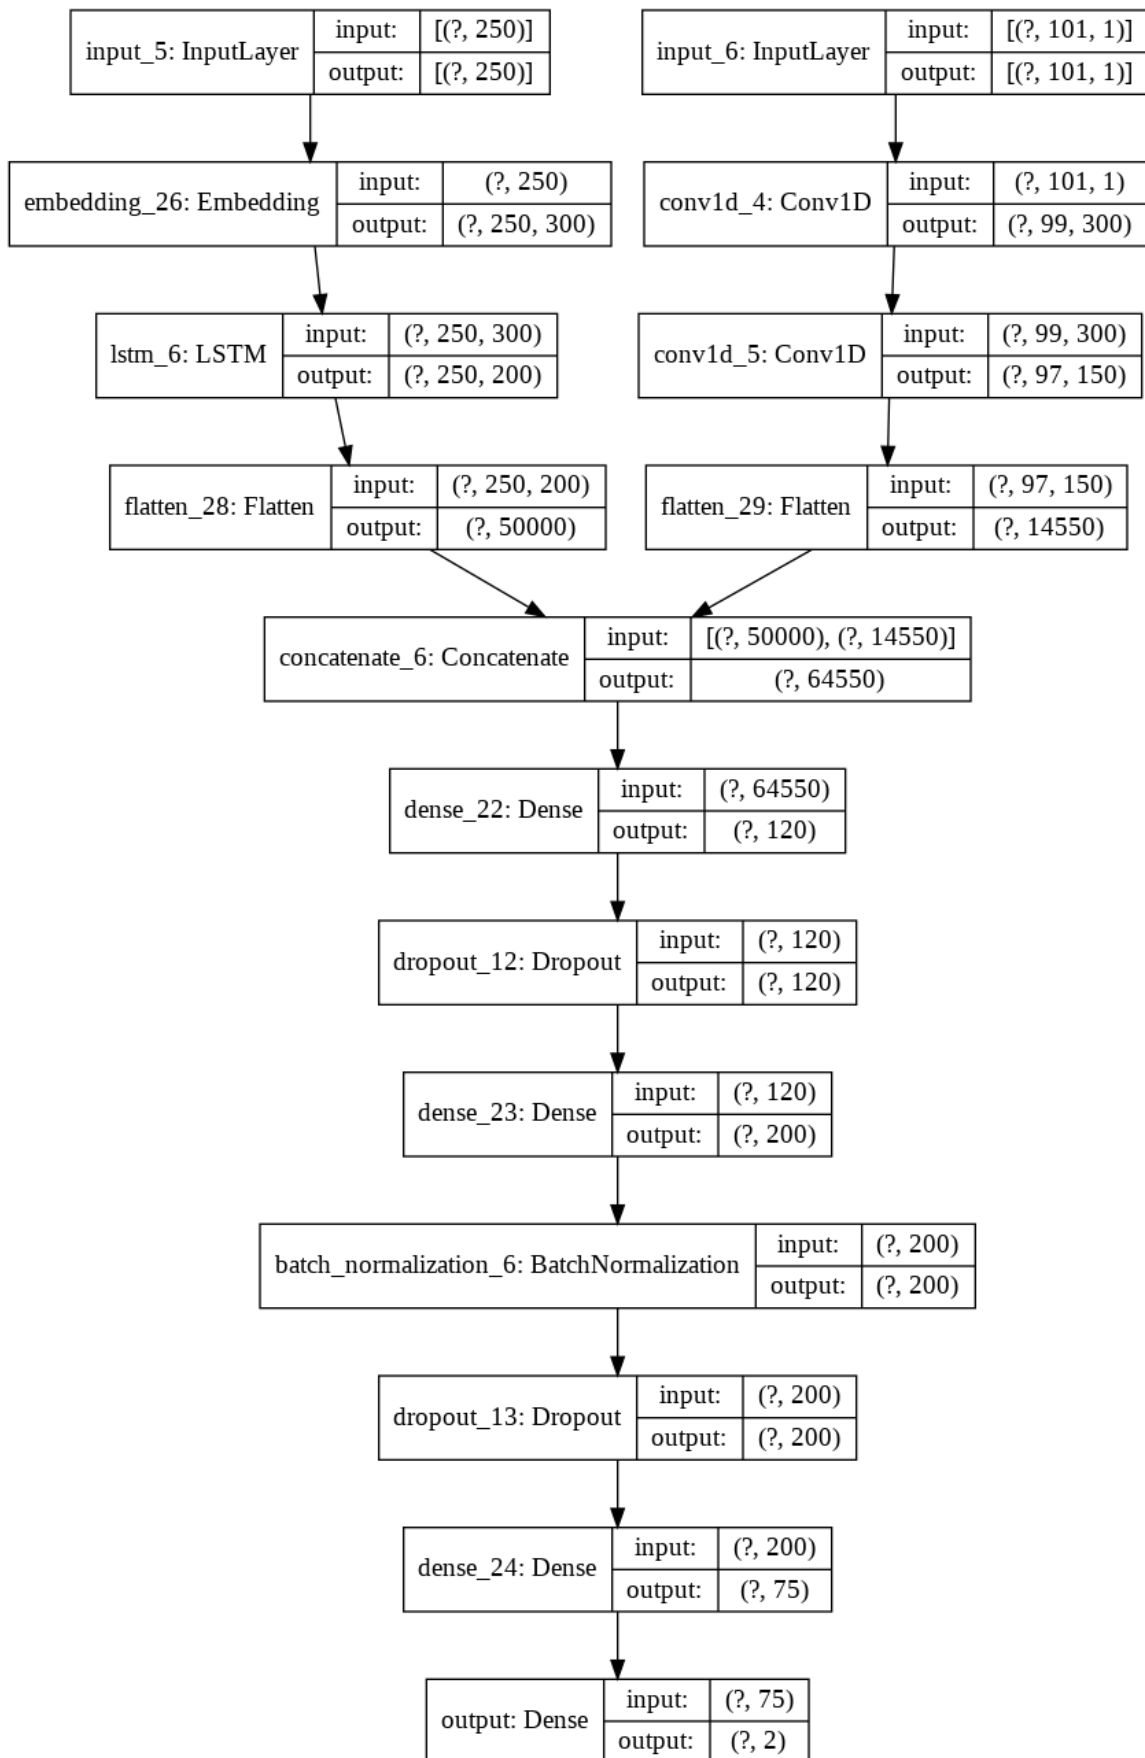
```
plt.plot(history2.history['auroc'], 'r')  
plt.plot(history2.history['val_auroc'], 'b')  
plt.legend({'Train ROCAUC': 'r', 'Test ROCAUC': 'b'})  
plt.show()
```



In [100]:

```
dot_img_file = '/tmp/model_3.png'  
tf.keras.utils.plot_model(model_3, to_file=dot_img_file, show_shapes=True)
```

Out[100]:



In []: